

Homework-1 Report

The aim of the homework is to implement `malloc()` and `free()` in C using linux BSD system call `sbrk()` which is a C wrapper on `brk()` system call.

The homework requires two different implementations of `malloc` - the first fit policy and best fit policy.

Design

The hardest part of this project was the design of the code. Various design decisions impact the performance and stability of the code. These were hard decisions to make.

Data Structure to track malloc/free data

The code uses a doubly linked list to track only the free or available memory segments. Having a doubly linked list was helpful in making modifications to the linked list, such as removing the identified node, adding a node, etc.

In order to be able to track total data segment size, I save the initial heap point as `heap_start` by saving the `sbrk(0)` value upon the first `malloc` call.

In order to track each memory location and its size, the linked list nodes are used as block metadata at the start of each malloced and free block. A malloced block contains a valid size but invalid pointers. A free block is a part of the linked list structure with valid pointers and size. When user gives a malloced pointer to free, the program adds the node to the linkedlist. This is then available to subsequent mallocs with its size already being valid. However, if the user frees a pointer not provided by `malloc`, the behaviour is undefined, similar to the functionality of the real `malloc`.

Malloc Calls

Whenever there is insufficient memory in the free list, an `sbrk()` system call is made. In order to improve the efficiency of this method, my program requests 4x the size requested by the user, including the metadata overhead. This reduces the number of system calls made in case of large mallocs with minimal frees.

Other Functions

insert_free() adds the node back to the free list

remove_free() removes the node from free list when malloced

Merge() and split() are required functions as well

For ff_malloc and bf_malloc, I use a common my_malloc function which takes an additional input *ff* to choose from ff_search and bf_search as the rest is identical between the two

Performance

First Fit performs better on small range allocation pattern as opposed to the large range allocation policies. The equal fit seems to show rather absurd results which may be due to some overflow possibility. Best Fit may not have much impact on equal range allocs but may improve performance for large range allocs.

First Fit Allocation Policy Performance Results

```
pr109@vcm-6252:~/ece650/Homework1/my_malloc/alloc_policy_tests$ ./large_range_rand_allocs
Execution Time = 181.139229 seconds
Fragmentation = 0.095369
```

```
pr109@vcm-6252:~/ece650/Homework1/my_malloc/alloc_policy_tests$ ./small_range_rand_allocs
data_segment_size = 3764256, data_segment_free_space = 261288
Execution Time = 10.984489 seconds
Fragmentation = 0.069413
```

```
pr109@vcm-6252:~/ece650/Homework1/my_malloc/alloc_policy_tests$ ./equal_size
Execution Time = 28.349003 seconds
Fragmentation = 2250.000000
pr109@vcm-6252:~/ece650/Homework1/my_malloc/alloc_policy_tests$ █
```